



C'est pas le réseau, ça
ping Marty!

Cynthia Treger
David Santiago

\$ ~~whoami~~ whoarewe



Cynthia Treger

Global Black Belt Azure Networking @Microsoft



[cynthiatreger](#)



David Santiago

Cloud Solution Architect @Microsoft



[davidsntg](#)

/ SUMMARY

1. Network 101

Connectivity, Routing, Filtering...

2. What could go wrong ?

Spoiler alert: It's not always DNS!

Network 101

/ Back to Basics


Networking example




IP addressing




Infrastructure



Infrastructure




Routing tables



Marty's Routing Table	
Route	Next Hop
<code>10.100.0.0/24</code>	<code>eth0</code>
<code>0.0.0.0/0</code>	<code>10.100.0.1</code>

R1's Routing Table	
Route	Next Hop
<code>10.100.0.0/24</code>	<code>eth0</code>
<code>192.168.11.0/24</code>	<code>eth1</code>

Routing tables

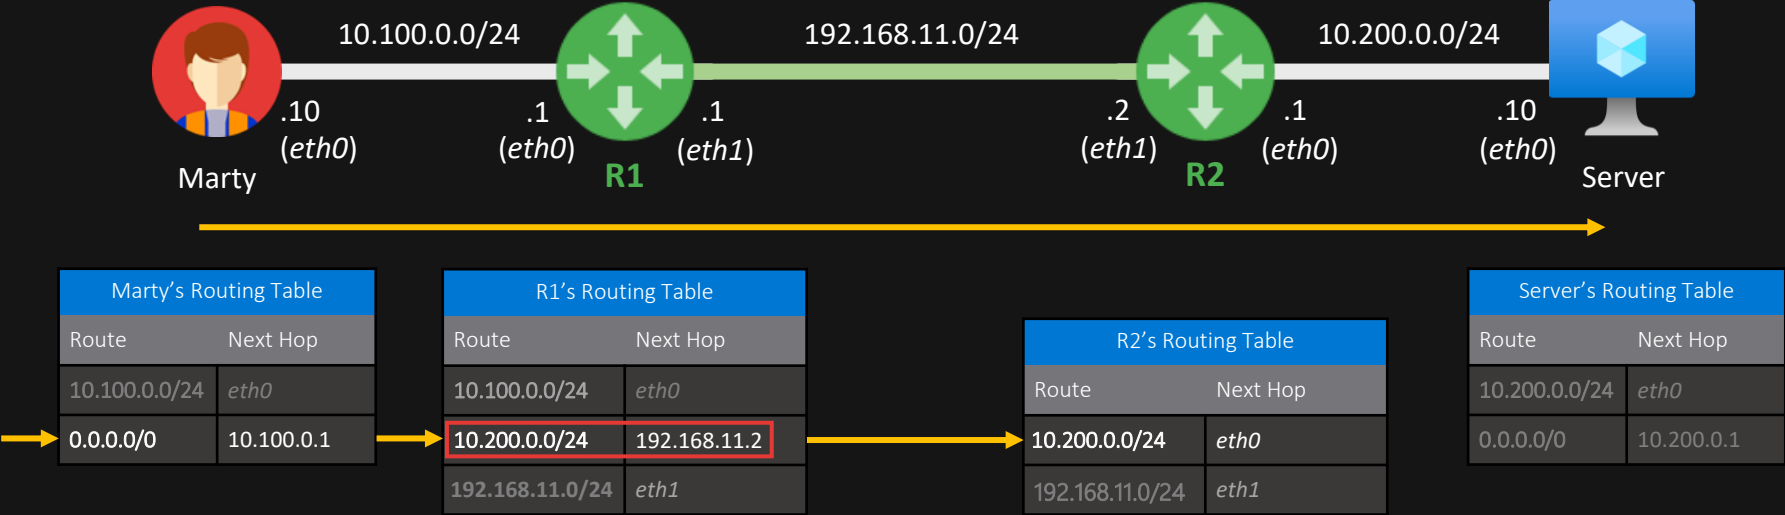


Marty's Routing Table	
Route	Next Hop
10.100.0.0/24	eth0
0.0.0.0/0	10.100.0.1


R1's Routing Table	
Route	Next Hop
10.100.0.0/24	eth0
192.168.11.0/24	eth1
10.200.0.0/24	192.168.11.2

Routing tables

TOOLS-IN-ACTION



Routing tables

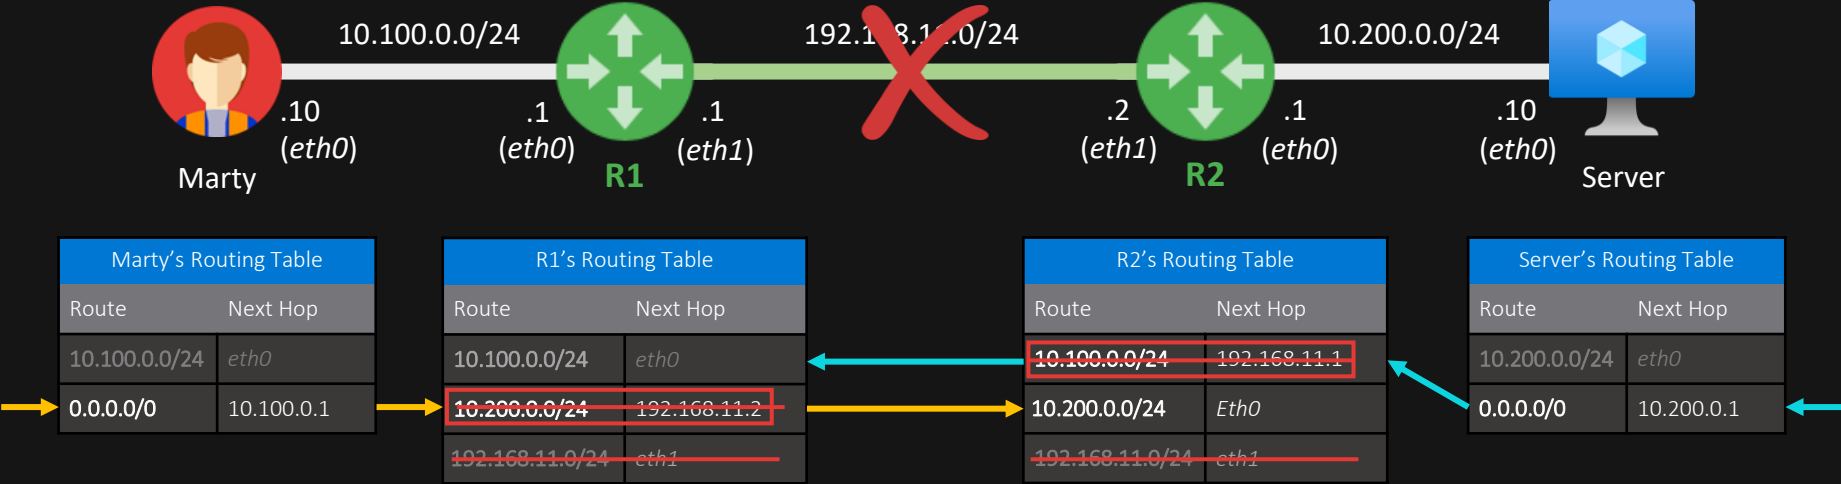


Marty's Routing Table		R1's Routing Table		R2's Routing Table		Server's Routing Table	
Route	Next Hop	Route	Next Hop	Route	Next Hop	Route	Next Hop
10.100.0.0/24	eth0	10.100.0.0/24	eth0	10.100.0.0/24	192.168.11.1	10.200.0.0/24	eth0
0.0.0.0/0	10.100.0.1	10.200.0.0/24	192.168.11.2	10.200.0.0/24	eth0	0.0.0.0/0	10.200.0.1
		192.168.11.0/24	eth1	192.168.11.0/24	eth1		

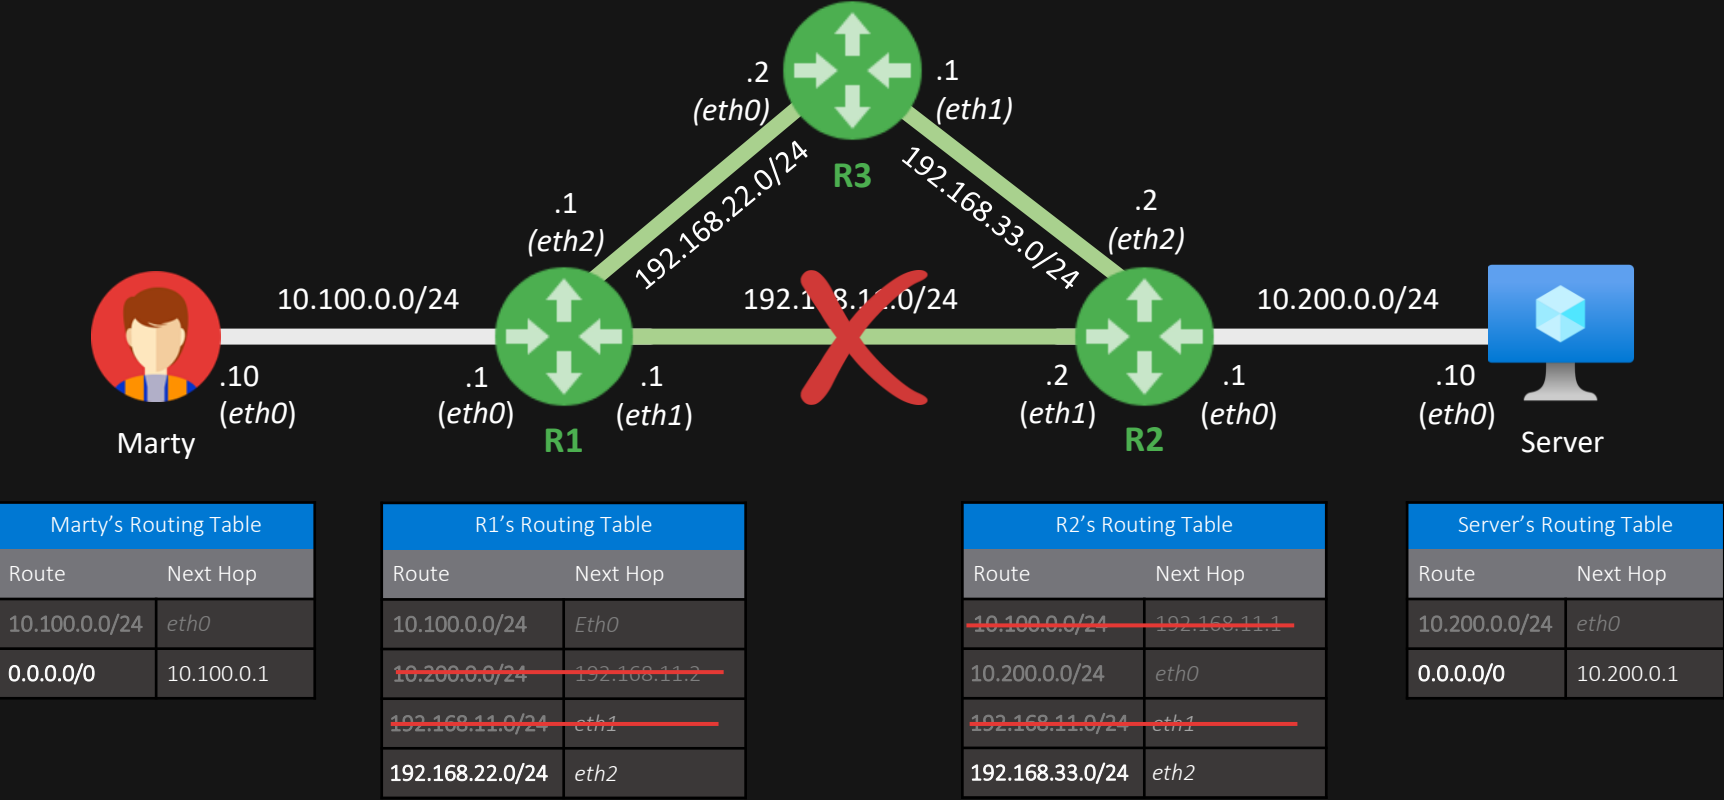


Civil works impact

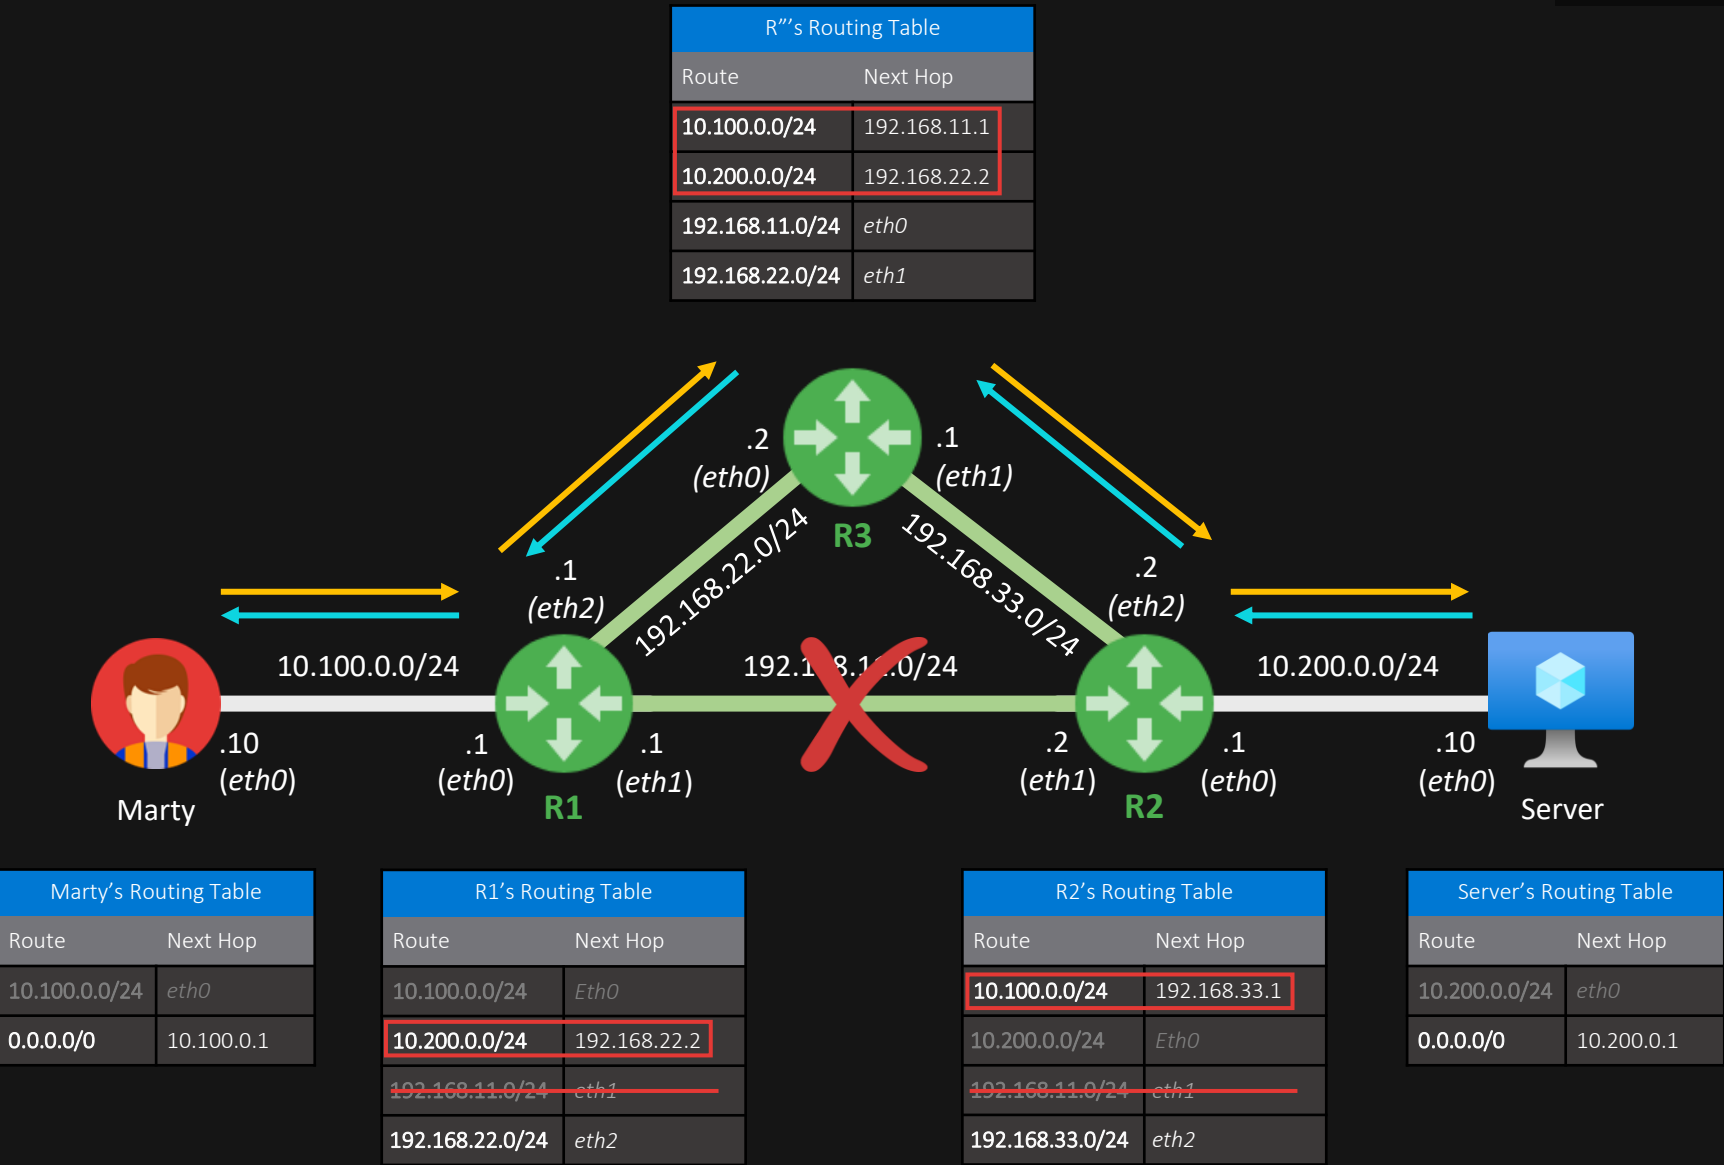
TOOLS-IN-ACTION



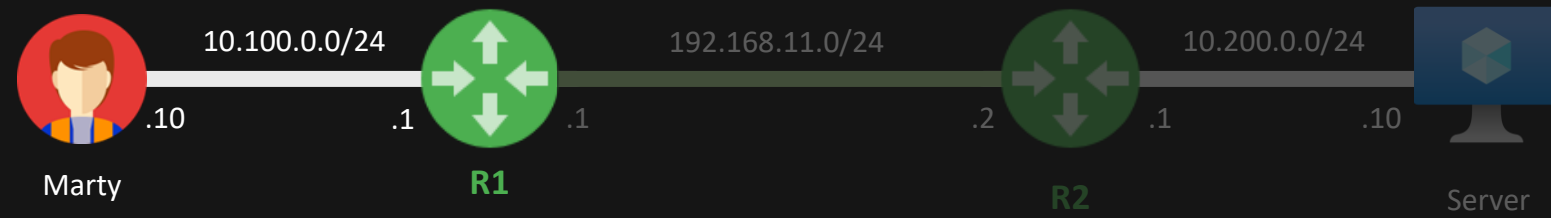
Path redundancy




Rerouting



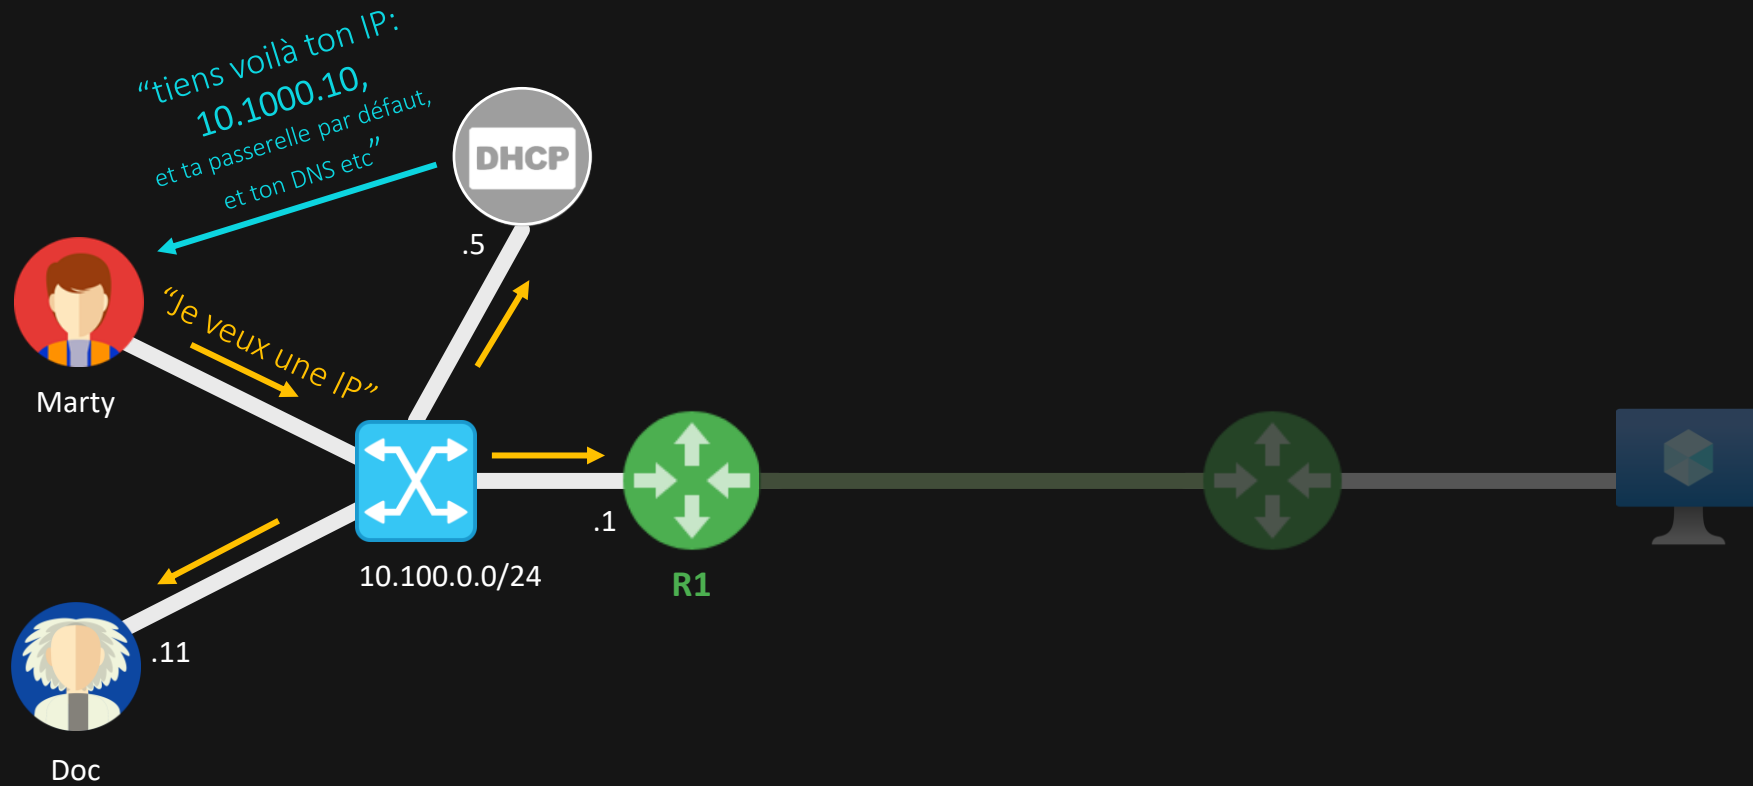
LAN connectivity




LAN connectivity



DHCP

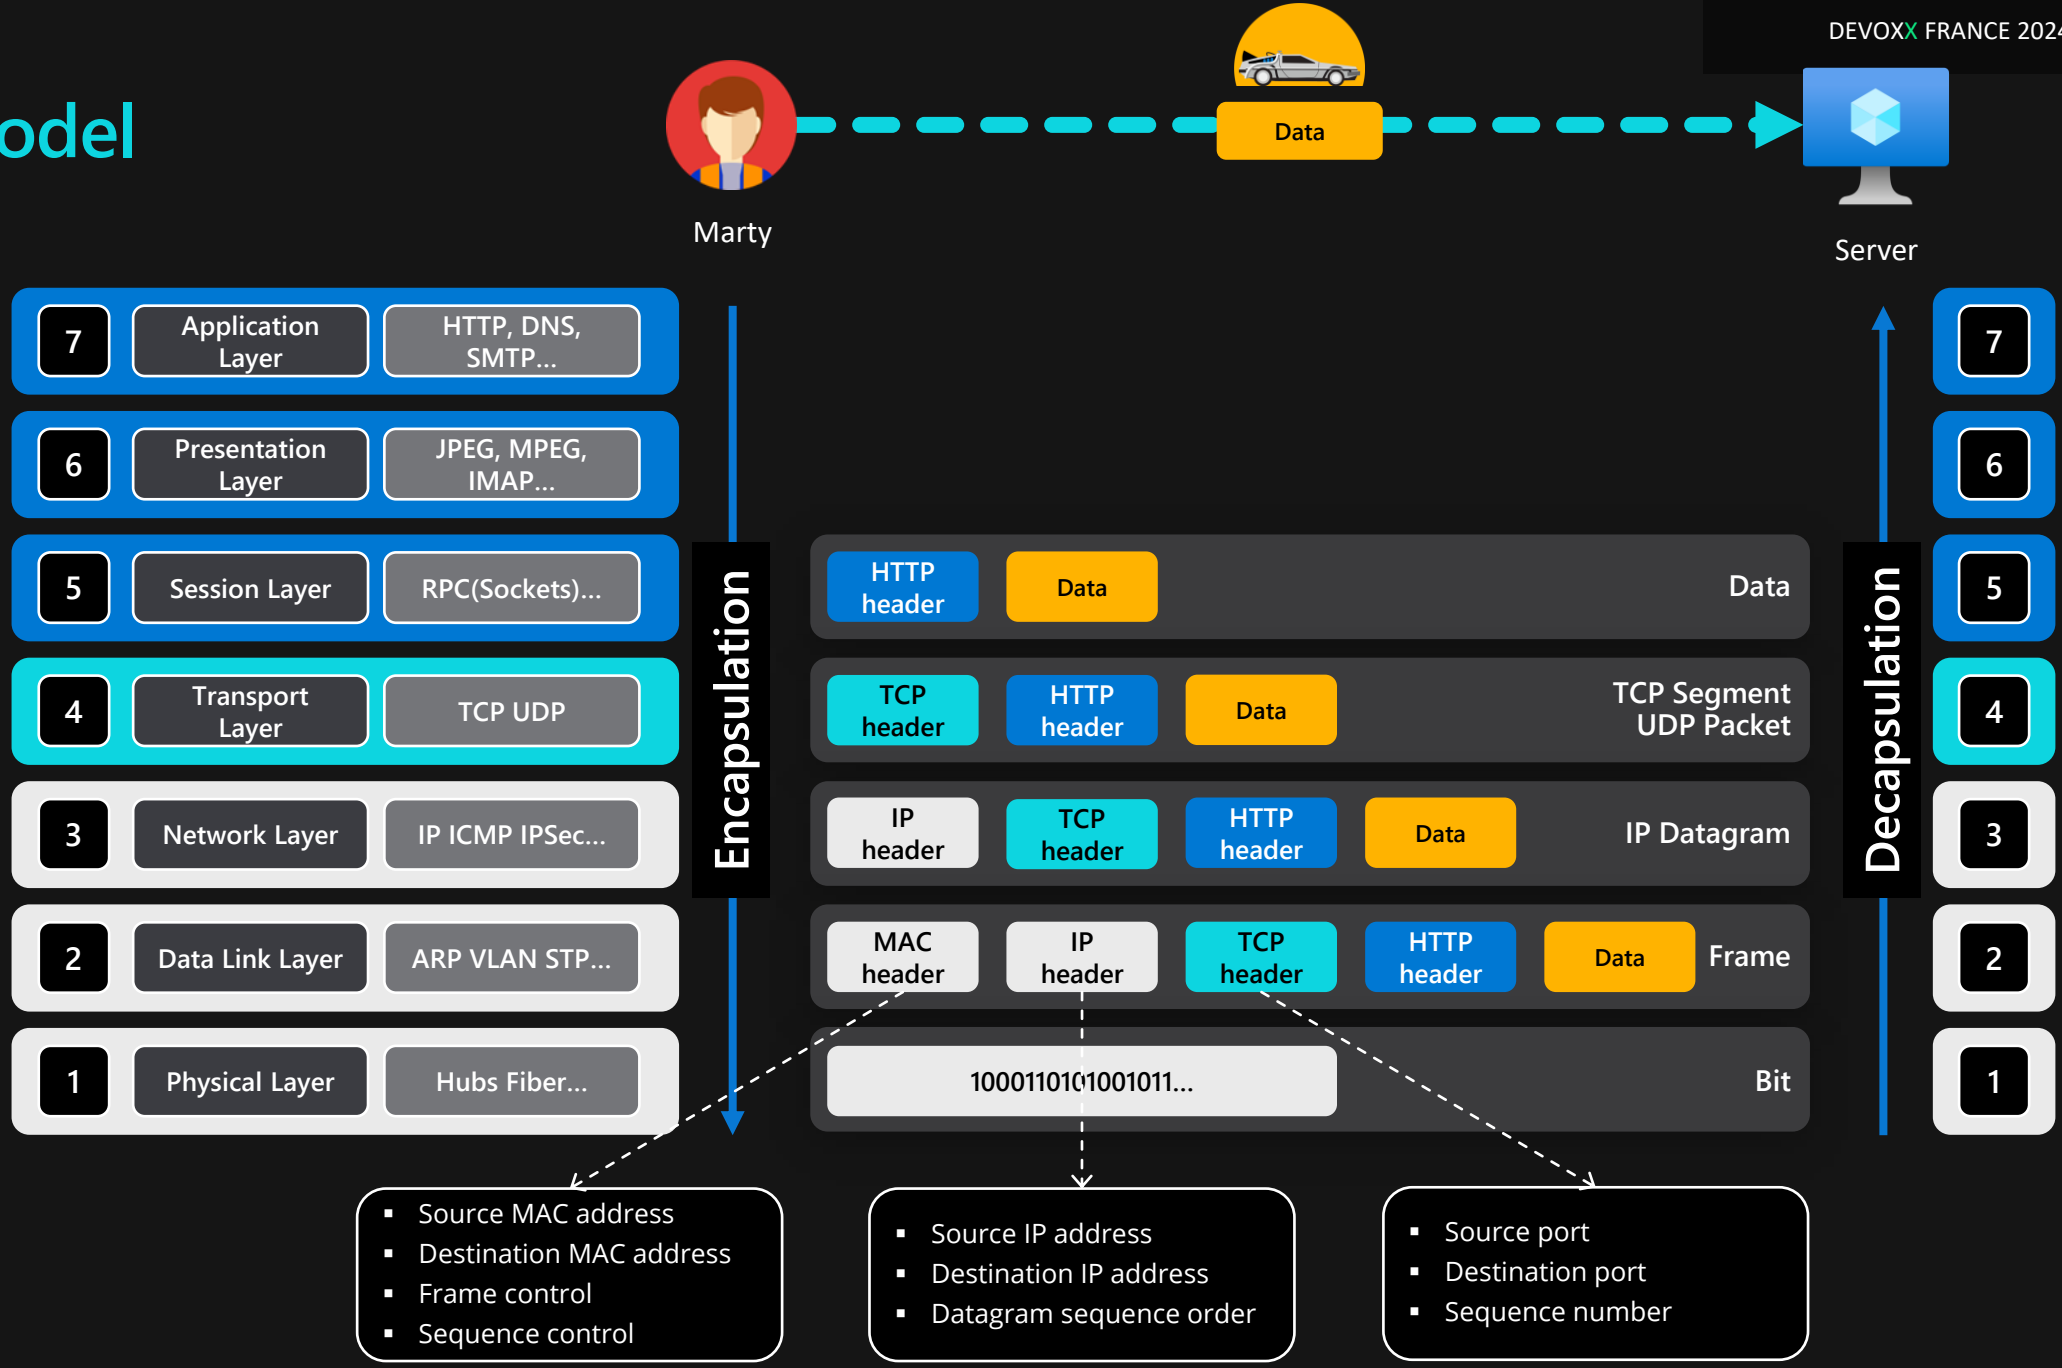


DNS

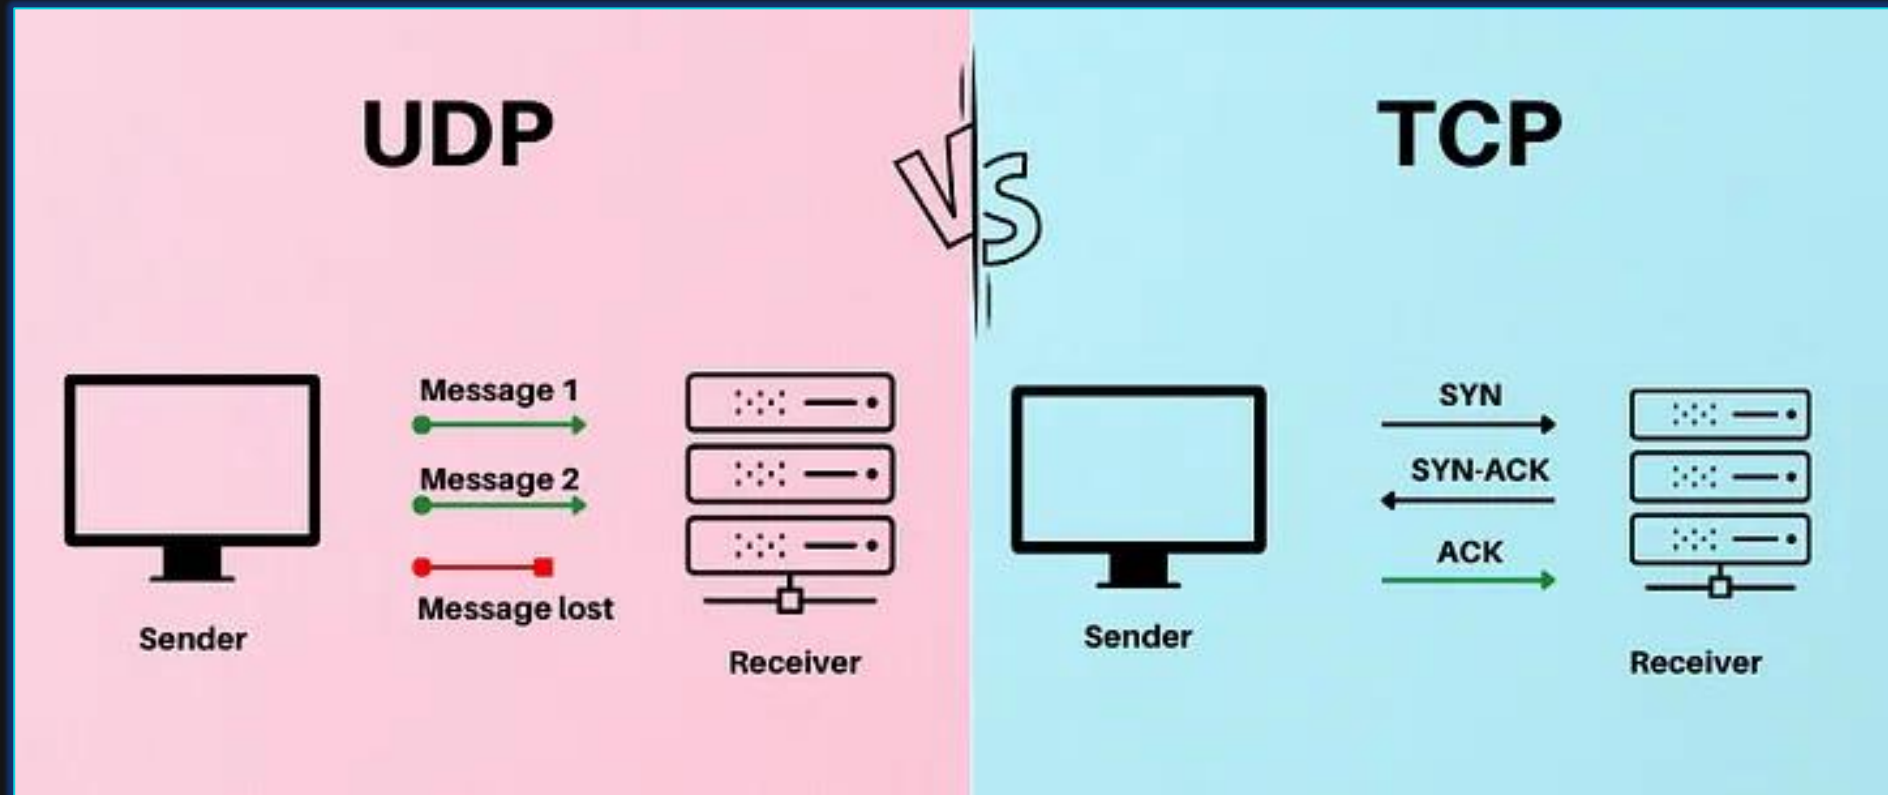


OSI model


TOOLS-IN-ACTION



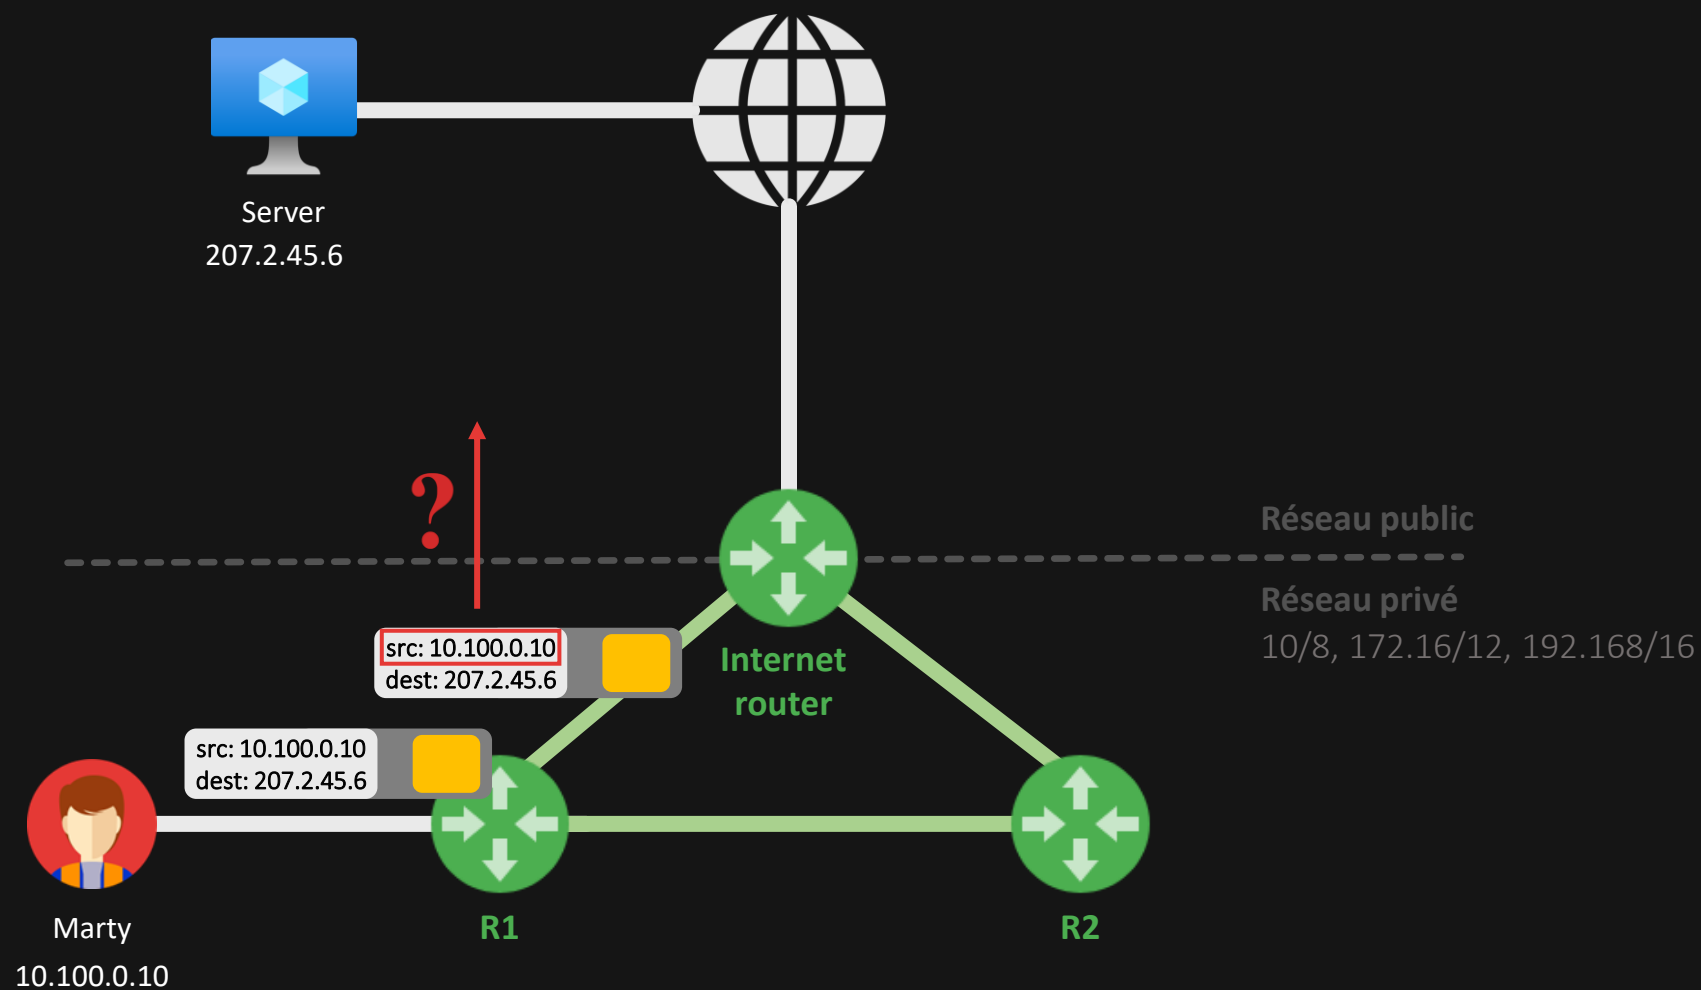
UDP vs TCP



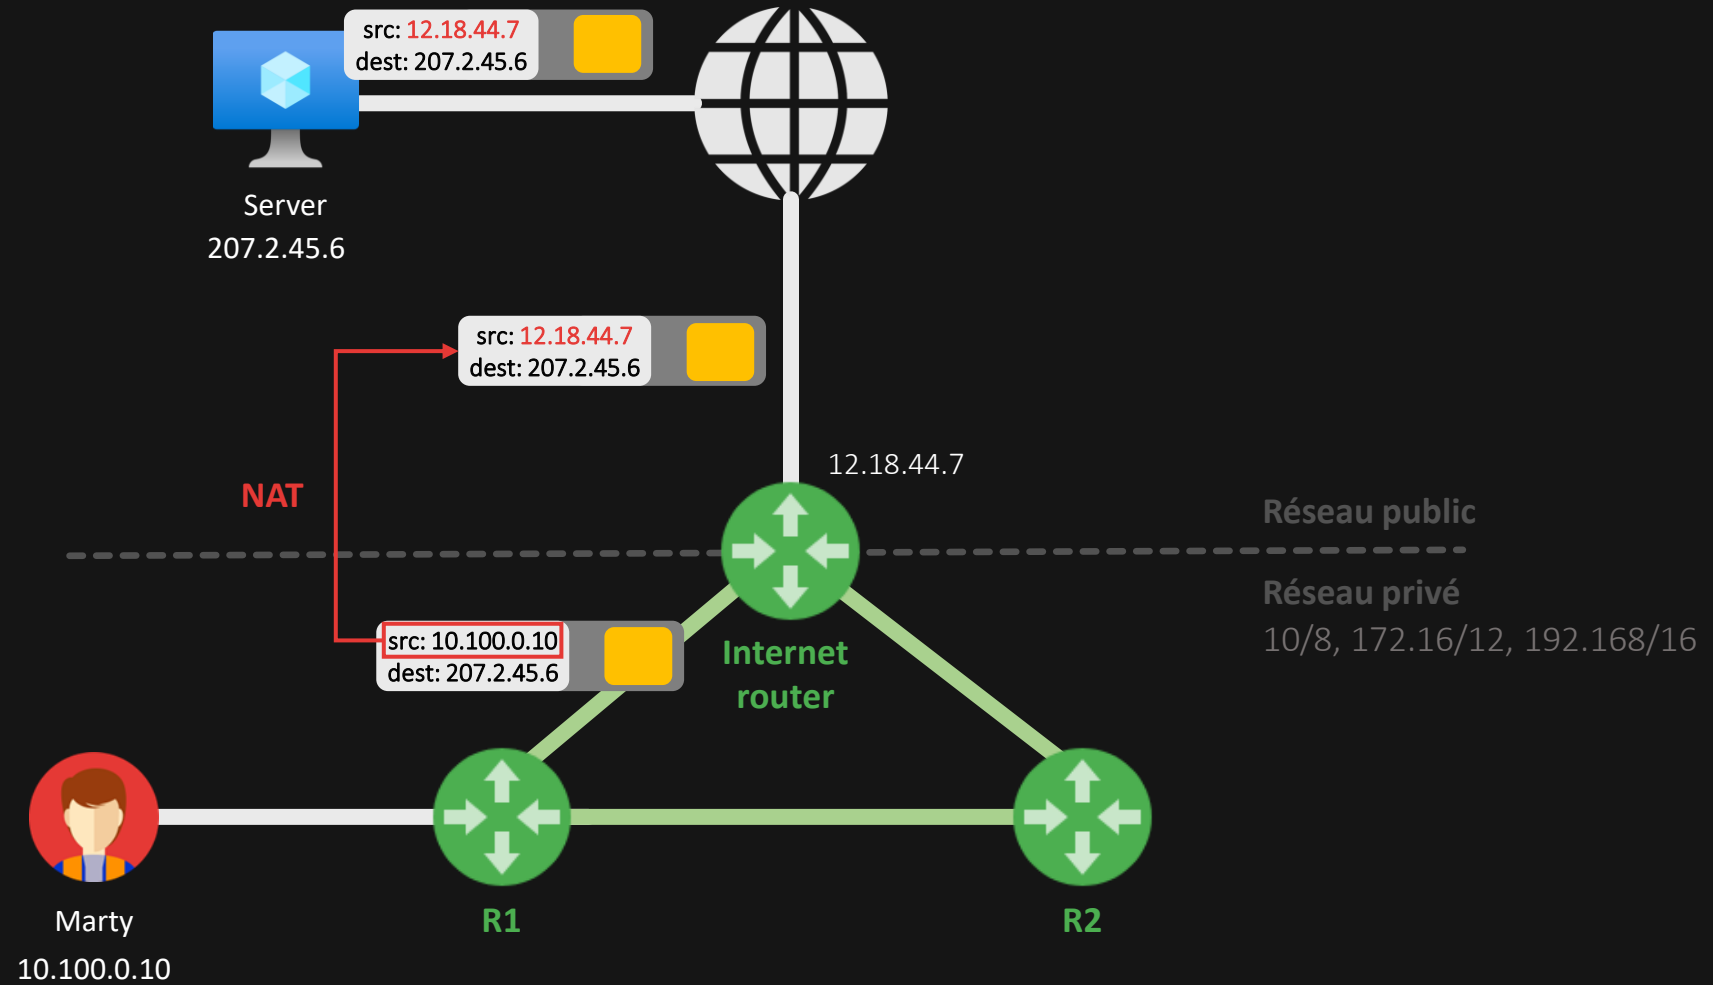
Private network vs public network



Internet access

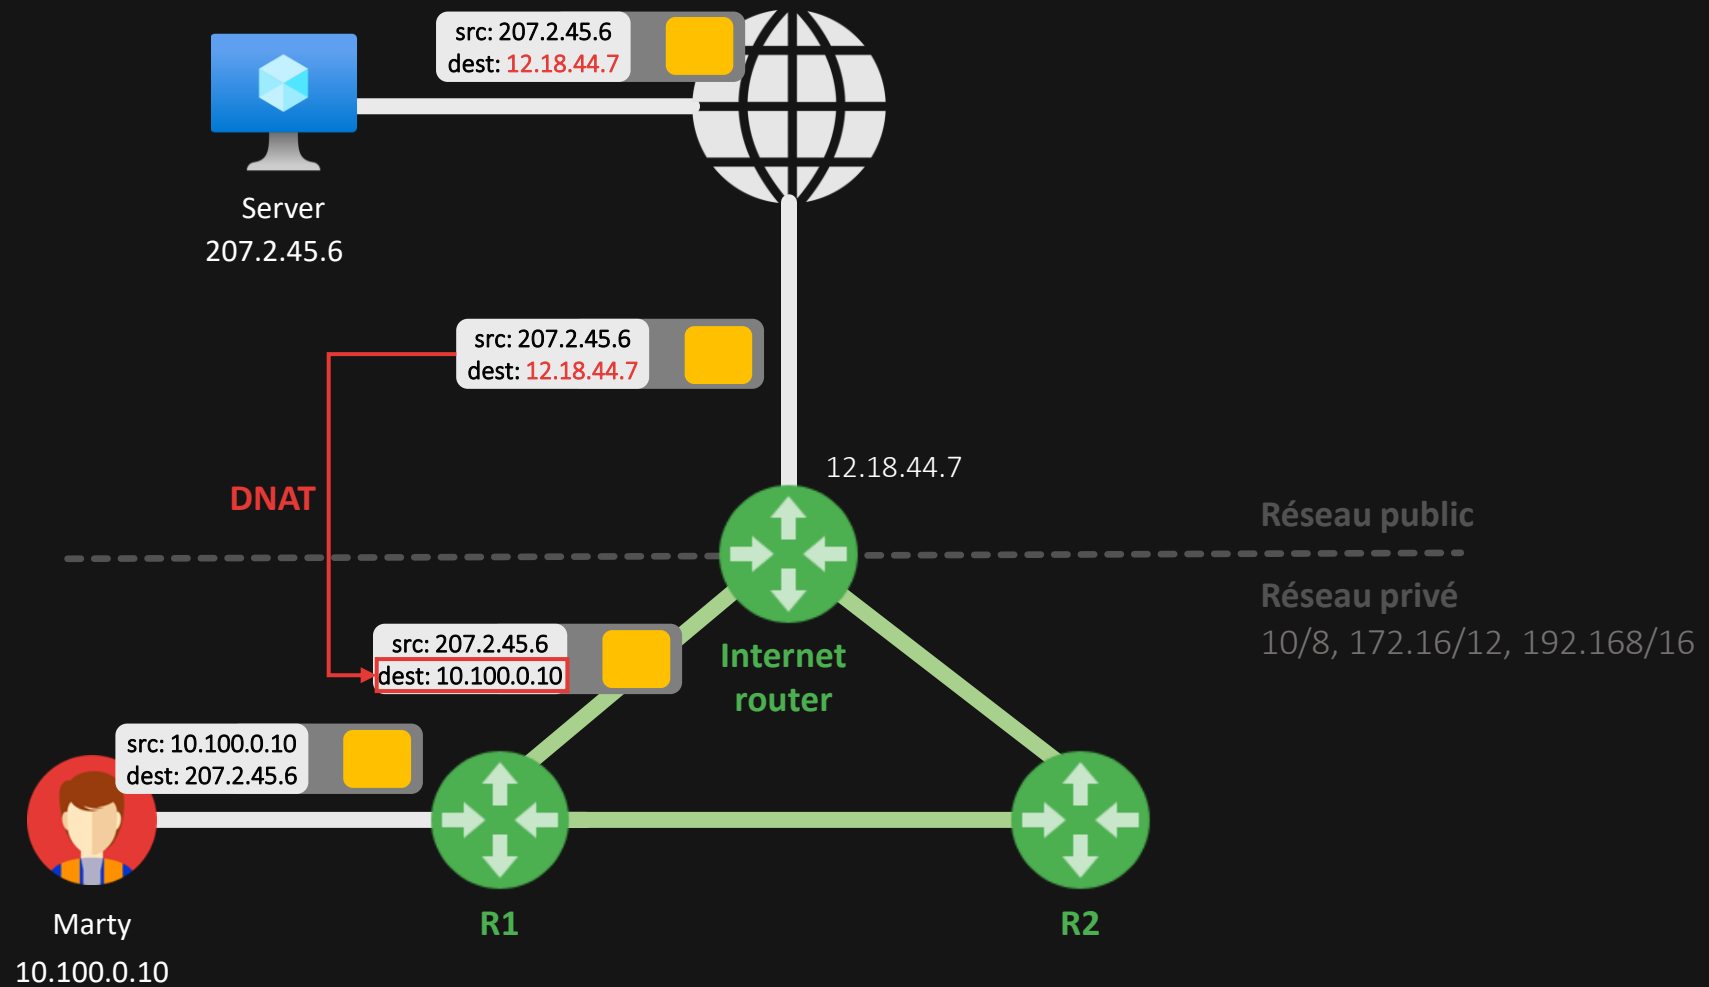


Internet access



Internet access

TOOLS-IN-ACTION



Latency, Bandwidth & Throughput

Latency

- The amount of time it takes data to move back and forth along a connection

Bandwidth

- The total capacity of a connection to move data

Throughput

- The actual amount of data that moves across a connection during a given time.

DALL·E Prompt: A close-up view of the Golden Gate Bridge in San Francisco, illuminated against the vibrant night sky, seen from an aerial perspective, reminiscent of a scene from a futuristic movie.



What could go wrong?

```
$ curl
```

- Command that lets you send or receive data from a website or file on a network.



```
$ curl 10.200.0.10
```


```
curl: (28) Failed to connect to 10.200.0.10 port 80: Connection timed out
```

\$ ping

How it works:

- Sends Internet Control Message Protocol (ICMP) packets to a specified address.
- If the target server returns an ICMP echo reply, it is connected to the network.

Example:



```
$ ping 10.200.0.10
PING 10.200.0.10 (10.200.0.10) 56(84) bytes of data.
^C
--- 10.200.0.10 ping statistics ---
3 packets transmitted, 0 received, 100% packet loss, time 2098ms
```

Possible Root Causes:

- The host is down
- There is no route from Marty to the host
- Firewall blocking the connection

Network troubleshooting #1

Display the kernel routing table

- Using `ip route`:



```
marty@devovx:~$ ip route
default via 10.100.0.1 dev eth0 proto dhcp src 10.100.0.10 metric 100
10.100.0.0/24 dev eth0 proto kernel scope link src 10.100.0.10
168.63.129.16 via 10.100.0.1 dev eth0 proto dhcp src 10.100.0.10 metric 100
169.254.169.254 via 10.100.0.1 dev eth0 proto dhcp src 10.100.0.10 metric 100
```

- Using `netstat`:



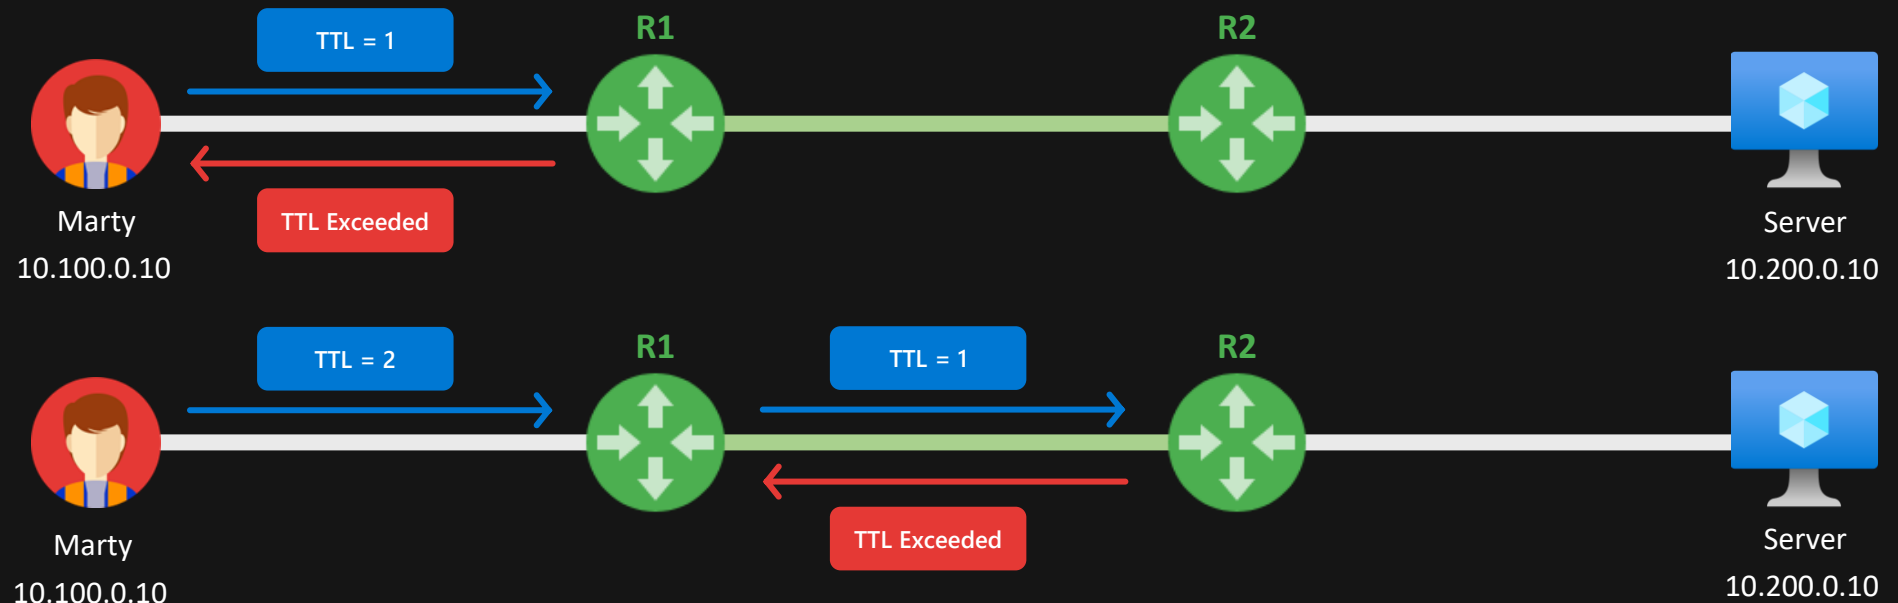
```
marty@devovx:~$ netstat -rn
Kernel IP routing table
Destination      Gateway           Genmask           Flags        MSS Window  irtt  Iface
0.0.0.0          10.100.0.1       0.0.0.0           UG            0 0        0     eth0
10.100.0.0       0.0.0.0          255.255.255.128  U            0 0        0     eth0
168.63.129.16    10.100.0.1       255.255.255.255  UGH           0 0        0     eth0
169.254.169.254  10.100.0.1       255.255.255.255  UGH           0 0        0     eth0
```

\$ traceroute


Two functions:

- Identifies each hop (that replies) in the path from a client computer to the destination entered in the `traceroute` command.
- Measures the round trip time for each hop in the path

How it works:





```
$ traceroute
```



```
marty@devoxx:~$ traceroute 10.200.0.10
traceroute to 10.200.0.10 (10.200.0.10), 30 hops max, 60 byte packets
 1  10.100.0.1 2.682 ms
 2  * * *
 3  * * *
 4  * * *
 5  * * *
```

```
$ traceroute
```



```
marty@devoxx:~$ traceroute 10.200.0.10
traceroute to 10.200.0.10 (10.200.0.10), 30 hops max, 60 byte packets
 1  10.100.0.1  2.920 ms
 2  192.168.11.2 (192.168.11.2)  5.085 ms
 3  10.200.0.10 (10.200.0.10)  6.980 ms
```

Network troubleshooting #2



```
marty@devoxx:~$ ping 10.200.0.10
PING 10.200.0.10 (10.200.0.10) 56(84) bytes of data.
64 bytes from 10.200.0.10: icmp_seq=1 ttl=63 time=6.74 ms
64 bytes from 10.200.0.10: icmp_seq=2 ttl=63 time=3.66 ms
64 bytes from 10.200.0.10: icmp_seq=3 ttl=63 time=3.88 ms
64 bytes from 10.200.0.10: icmp_seq=4 ttl=63 time=5.21 ms
^C
--- 10.200.0.10 ping statistics ---
4 packets transmitted, 4 received, 0% packet loss, time 3004ms
rtt min/avg/max/mdev = 3.659/4.872/6.744/1.231 ms
```



```
marty@devoxx:~$ curl 10.200.0.4
curl: (7) Failed to connect to 10.200.0.4 port 80: Connection refused
```

Possible Root Causes:

- No service listening on the specified port (:80)
- Firewall blocking the connection

\$ netstat

- Command that show network connections, routing tables, and a number of network interface & network protocol statics.




```
root@server:~$ netstat -tulnp
```

```
Active Internet connections (only servers)
```

Proto	Recv-Q	Send-Q	Local Address	Foreign Address	State	PID/Program name
tcp	0	0	127.0.0.53:53	0.0.0.0:*	LISTEN	575/systemd-resolve
tcp	0	0	0.0.0.0:22	0.0.0.0:*	LISTEN	857/sshd: /usr/sbin
tcp6	0	0	:::22	:::*	LISTEN	857/sshd: /usr/sbin
udp	0	0	127.0.0.1:323	0.0.0.0:*		714/chronyd
udp	0	0	127.0.0.1:25224	0.0.0.0:*		1135/ruby
udp	0	0	127.0.0.53:53	0.0.0.0:*		575/systemd-resolve
udp	0	0	10.100.0.4:68	0.0.0.0:*		573/systemd-network
udp6	0	0	:::1:323	:::*		714/chronyd

```
$ curl
```



```
marty@devoxx:~$ curl 10.200.0.10
```

```
Hello Devoxx France! 🚚
```

\$ sockperf – *for latency, packet loss ...*

```

marty@devovx:~$ sockperf pp -i 10.200.0.10 -m 64

sockperf: == version #3.6-no.git ==
sockperf[CLIENT] send on:sockperf: using recvfrom() to block on socket(s)

[ 0] IP = 10.200.0.10      PORT = 11111 # UDP
sockperf: Warmup stage (sending a few dummy messages)...
sockperf: Starting test...
sockperf: Test end (interrupted by timer)
sockperf: Test ended
sockperf: [Total Run] RunTime=1.000 sec; Warm up time=400 msec; SentMessages=192; ReceivedMessages=191
sockperf: ===== Printing statistics for Server No: 0
sockperf: [Valid Duration] RunTime=0.552 sec; SentMessages=107; ReceivedMessages=107
sockperf: ==> avg-latency=2579.605 (std-dev=396.137)
sockperf: # dropped messages = 0; # duplicated messages = 0; # out-of-order messages = 0
sockperf: Summary: Latency is 2579.605 usec
sockperf: Total 107 observations; each percentile contains 1.07 observations
sockperf: ---> <MAX> observation = 5068.870
sockperf: ---> percentile 99.999 = 5068.870
sockperf: ---> percentile 99.990 = 5068.870
sockperf: ---> percentile 99.900 = 5068.870
sockperf: ---> percentile 99.000 = 3932.427
sockperf: ---> percentile 90.000 = 2911.623
sockperf: ---> percentile 75.000 = 2690.890
sockperf: ---> percentile 50.000 = 2474.818
sockperf: ---> percentile 25.000 = 2339.955
sockperf: ---> <MIN> observation = 2016.383

```

\$ sockperf – *for throughput*



```
marty@devovx:~$ sockperf throughput --tcp -i 10.200.0.10 -p 12345 -m 1500 -t 20
sockperf: == version #3.6-no.git ==
sockperf[CLIENT] send on:
[ 0] IP = 10.200.0.10      PORT = 12345 # TCP
sockperf: Warmup stage (sending a few dummy messages)...
sockperf: Starting test...
sockperf: Test end (interrupted by timer)
sockperf: Test ended
sockperf: Total of 1535781 messages sent in 20.001 sec

sockperf: NOTE: test was performed, using msg-size=1500. For getting maximum throughput consider using --msg-size=1472
sockperf: Summary: Message Rate is 76787 [msg/sec], Packet Rate is about 153574 [pkt/sec] (2 ip frags / msg)
sockperf: Summary: BandWidth is 109.845 MBps (878.757 Mbps)
```

Alternatives:

- Ntttcp, iperf, iperf3, Tcptrack

\$ dig – *for DNS*

```
marty@devovx:~$ dig www.backtothefuture.com

; <<>> DiG 9.16.48-Ubuntu <<>> www.backtothefuture.com
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 10964
;; flags: qr rd ra; QUERY: 1, ANSWER: 5, AUTHORITY: 0, ADDITIONAL: 1

;; OPT PSEUDOSECTION:
; EDNS: version: 0, flags:; udp: 65494
;; QUESTION SECTION:
;www.backtothefuture.com.      IN      A

;; ANSWER SECTION:
www.backtothefuture.com. 1800    IN      CNAME   ext-sq.squarespace.com.
ext-sq.squarespace.com. 6       IN      A       198.49.23.145
ext-sq.squarespace.com. 6       IN      A       198.185.159.144
ext-sq.squarespace.com. 6       IN      A       198.49.23.144
ext-sq.squarespace.com. 6       IN      A       198.185.159.145

;; Query time: 32 msec
;; SERVER: 127.0.0.53#53(127.0.0.53)
;; WHEN: Thu Apr 04 10:37:08 UTC 2024
;; MSG SIZE rcvd: 149
```




THANKS
FOR WATCHING